



KTU  
**NOTES**  
The learning companion.

**KTU STUDY MATERIALS | SYLLABUS | LIVE  
NOTIFICATIONS | SOLVED QUESTION PAPERS**

🌐 Website: [www.ktunotes.in](http://www.ktunotes.in)



# CST 402 - DISTRIBUTED COMPUTING

Recitations: M

Ktunotes.in  
**Module - III**

# Module – III

## Lesson Plan

- **L1:** Distributed mutual exclusion algorithms – System model, Lamport's algorithm
- **L2:** Ricart–Agrawala algorithm
- **L3:** Quorum-based mutual exclusion algorithms – Maekawa's algorithm
- **L4:** Token-based algorithm – Suzuki–Kasami's broadcast algorithm.
- **L5:** Deadlock detection in distributed systems – System model, Deadlock handling strategies, Issues in deadlock detection.
- **L6:** Models of deadlocks

# Distributed mutual exclusion algorithms

- Mutual exclusion is a fundamental problem in distributed computing systems.
- Mutual exclusion ensures that concurrent access of processes to a shared resource or data is serialized, that is, executed in a mutually exclusive manner.
- Mutual exclusion in a distributed system states that only one process is allowed to execute the critical section (CS) at any given time
- There are for implementing distributed mutual exclusion:
  1. Token-based approach.
  2. Non-token-based approach
  3. Quorum-based approach.

# Distributed mutual exclusion algorithms

- **In the token-based approach**, a unique token is shared among the sites.
- A site is allowed to enter its CS if it possesses the token and it continues to hold the token until the execution of the CS is over.
- Mutual exclusion is ensured because the token is unique
- **In the non-token-based approach**, two or more successive rounds of **selection** among the sites to determine which site will enter the CS next.
- A site enters the critical section (CS) when an assertion, defined on its local variables, becomes true.
- **In the quorum-based approach**, each site requests permission to execute the CS from a subset of sites (called a quorum).
- The quorums are formed in such a way that when two sites concurrently request access to the CS, at least one site receives both the requests and this site is responsible to make sure that only one request executes the CS at any time.

# Distributed mutual exclusion algorithms

- **System model**
- The system consists of **N sites**, S<sub>1</sub>, S<sub>2</sub>, , S<sub>N</sub>. Without loss of generality, we assume that a single process is running on each site
- The process at site S<sub>i</sub> is denoted by p<sub>i</sub>.
- All these processes communicate over an underlying communication network.
- A process wishing to enter the CS sends REQUEST messages to all other or a subset of processes by sending REQUEST messages, and waits for appropriate responses before entering the CS.
- A site can be in one of the following states: requesting the CS, executing the CS, or neither requesting nor executing the CS.
- A site can be in one of the following states: requesting the CS, executing the CS, or neither requesting nor executing the CS.

## Distributed mutual exclusion algorithms

- In the “requesting the CS” state, the site is blocked and cannot make further requests for the CS.
  - In the idle state, the site is executing outside the CS.
  - In the token-based algorithms, a site can also be in a state idle token
- Ktunotes.in**
- Such state is referred to as the idle token state.
  - At any instant, a site may have several pending requests for CS.
  - A site receives requests and serves them one at a time.

# Distributed mutual exclusion algorithms

- We do not make any assumption regarding communication channels if they are FIFO or not.
- This [redacted] [redacted]

## Requirements of mutual exclusion algorithms

A mutual exclusion algorithm should satisfy the following properties:

### 1. Safety property:

The safety property states that at any instant, [redacted] section.

This is an essential property of a mutual exclusion algorithm.

### 2. Liveness property :

This property states th [redacted]

# Distributed mutual exclusion algorithms

Two or more sites should not endlessly wait for messages that will never arrive.

In addition, a site must not wait indefinitely to execute the CS while other sites are repeatedly executing the CS.

That is,

Ktunotes.in

## 3. Fairness :

Fairness in the context of mutual exclusion means that each process gets a chance to execute the CS.

In mutual exclusion algorithms, the fairness property generally means that the CS execution requests are executed in a round-robin fashion.

# Distributed mutual exclusion algorithms

## Performance metrics

The performance of mutual exclusion algorithms is generally measured by the following

**Message complexity** : This is the number of messages that are required per CS execution by a site

Ktunotes.in

**Synchronization delay** : After a site leaves the CS, it is the time required and before the next site enters the CS

**Response time** : This is the time interval a request waits for its CS execution to be over after its request messages have been sent out

# Distributed mutual exclusion algorithms

## Performance metrics

**System throughput** This is the rate at which the system executes requests for the CS. If  $SD$  is the synchronization delay and  $E$  is the average critical section execution time, then the throughput is given by the following equation:

$$\text{System throughput} = \frac{1}{(SD + E)}.$$

Ktunotes.in

## Lamport's algorithm

Lamport developed a distributed mutual exclusion algorithm

The algorithm is fair in the sense that a request for CS are executed in the order of their timestamps and t

it updates its local clock and assigns the request a timestamp.

The algorithm executes CS requests i

Every site  $S_i$  keeps a queue, which contains mutual exclusion requests ordered by their timestamps.

This algorithm requires c deliver messages in order.

# Lamport's algorithm

---

## Requesting the critical section

- When a site  $S_i$  wants to enter the CS, it broadcasts a REQUEST( $ts_i, i$ ) message to all other sites and places the request on  $request\_queue_i$ . (( $ts_i, i$ ) denotes the timestamp of the request.)
- When a site  $S_j$  receives the REQUEST( $ts_i, i$ ) message from site  $S_i$ , it places site  $S_i$ 's request on  $request\_queue_j$  and returns a timestamped REPLY message to  $S_i$ .

## Executing the critical section

Site  $S_i$  enters the CS when the following two conditions hold:

**L1:**  $S_i$  has received a message with timestamp larger than  $(ts_i, i)$  from all other sites.

**L2:**  $S_i$ 's request is at the top of  $request\_queue_i$ .

## Releasing the critical section

- Site  $S_i$ , upon exiting the CS, removes its request from the top of its request queue and broadcasts a timestamped RELEASE message to all other sites.
- When a site  $S_j$  receives a RELEASE message from site  $S_i$ , it removes  $S_i$ 's request from its request queue.

---

**Algorithm 9.1** Lamport's algorithm.

## Ricart–Agrawala algorithm

- The Ricart–Agrawala algorithm assumes that the communication channels are FIFO.
- The algorithm uses two types of messages: REQUEST and REPLY.
- A process sends a REQUEST message to all other processes to request their permission to enter the critical section.
- A process sends a REPLY message to a process to give its permission to that process.
- Processes use timestamps to assign a timestamp to critical section requests.
- Timestamps are used to decide the priority of requests in case of conflict

## Ricart–Agrawala algorithm

- if a process  $p_i$  that is waiting to execute the critical section receives a REQUEST message from process  $p_j$ ,
- then if the priority of  $p_j$ 's request is lower,  $p_i$  defers the REPLY to  $p_j$  and sends a REPLY message to  $p_j$  only after executing the CS for its pending request.
- Otherwise,  $p_i$  sends a REPLY message to  $p_j$  immediately, provided it is currently not executing the CS.
- Each process  $p_i$  maintains the pending requests queue, the size of which is the same as the number of processes in the system.

# Ricart–Agrawala algorithm

---

## Requesting the critical section

- (a) When a site  $S_i$  wants to enter the CS, it broadcasts a timestamped REQUEST message to all other sites.
- (b) When site  $S_j$  receives a REQUEST message from site  $S_i$ , it sends a REPLY message to site  $S_i$  if site  $S_j$  is neither requesting nor executing the CS, or if the site  $S_j$  is requesting and  $S_i$ 's request's timestamp is smaller than site  $S_j$ 's own request's timestamp. Otherwise, the reply is deferred and  $S_j$  sets  $RD_j[i] := 1$ .

## Executing the critical section

- (c) Site  $S_i$  enters the CS after it has received a REPLY message from every site it sent a REQUEST message to.

## Releasing the critical section

- (d) When site  $S_i$  exits the CS, it sends all the deferred REPLY messages;  $\forall j$  if  $RD_i[j] = 1$ , then sends a REPLY message to  $S_j$  and sets  $RD_i[j] := 0$ .

---

**Algorithm 9.2** The Ricart–Agrawala algorithm.

## Quorum-based mutual exclusion algorithms

- Quorum-based mutual exclusion algorithms represented a departure from the trend in the following two ways:
- A site does not request permission from all other sites, but only from a subset of the sites.
- This is a radically different approach as compared to the Lamport and Ricart–Agrawala algorithms,
- In quorum-based mutual exclusion algorithm,
- A site can send a REPLY message only after it has received a RELEASE message for the previous REPLY message.
- Therefore, a site  $S_i$  locks all the sites in  $R_i$  in exclusive mode before executing its CS.

# Quorum-based mutual exclusion algorithms

- Quorum-based mutual exclusion algorithms significantly reduce the message complexity of invoking mutual exclusion by having sites ask permission from only a subset of sites.
- Since these algorithms are based on the notion of “Coteries” and “Quorums,” we first describe the idea of coteries and quorums.
- The following properties
- I
- Coterie
- Quorum
- Coteries and quorums can be used to develop algorithms to ensure mutual exclusion in a distributed environment

## Quorum-based mutual exclusion algorithms

- A simple protocol works as follows: let “a” be a site in quorum “A.”
- If “a” wants to invoke mutual exclusion, it requests permission from all sites in its quorum “A.”
- Minimality property ensures efficiency

Ktunotes.in

# Maekawa's algorithm

- Maekawa's algorithm was the first quorum-based mutual exclusion algorithm.
  - This algorithm requires delivery of messages to be guaranteed before they are sent between every pair of sites.
- 

## Requesting the critical section:

- (a) A site  $S_i$  requests access to the CS by sending REQUEST( $i$ ) messages to all sites in its request set  $R_i$ .
- (b) When a site  $S_j$  receives the REQUEST( $i$ ) message, it sends a REPLY( $j$ ) message to  $S_i$  provided it hasn't sent a REPLY message to a site since its receipt of the last RELEASE message. Otherwise, it queues up the REQUEST( $i$ ) for later consideration.

## Executing the critical section:

- (c) Site  $S_i$  executes the CS only after it has received a REPLY message from every site in  $R_i$ .

## Releasing the critical section:

- (d) After the execution of the CS is over, site  $S_i$  sends a RELEASE( $i$ ) message to every site in  $R_i$ .
- (e) When a site  $S_j$  receives a RELEASE( $i$ ) message from site  $S_i$ , it sends a REPLY message to the next site waiting in the queue and deletes that entry from the queue. If the queue is empty, then the site updates its state to reflect that it has not sent out any REPLY message since the receipt of the last RELEASE message.

---

## Algorithm 9.5 Maekawa's algorithm.

## Token-based algorithms

- In token-based algorithms, a unique token is shared among the sites.
- A site is allowed to enter its CS if it possesses the token.
- A site holding the token can enter its CS repeatedly until it sends the token to some other site.

- There are numerous token-based algorithms
- token-based algorithms instead of timestamps.
-

## Suzuki–Kasami's broadcast algorithm

- In Suzuki–Kasami's algorithm if a site that wants to enter the CS does not have the token, it broadcasts a REQUEST message for the token to all other sites.
- A site that possesses the token sends it to the requesting site upon the receipt of its REQUEST message.
- If a site receives a REQUEST message it sends the token only after it has completed the execution of the CS
- Although the basic idea underlying this algorithm may sound rather simple, there are that must be efficiently addressed:
  - 1. How to distinguishing an outdated REQUEST message from a current REQUEST message
  - 2. How to determine which site has an outstanding request for the CS

# Suzuki–Kasami’s broadcast algorithm

---

## Requesting the critical section:

- (a) If requesting site  $S_i$  does not have the token, then it increments its sequence number,  $RN_i[i]$ , and sends a REQUEST( $i, sn$ ) message to all other sites. (“ $sn$ ” is the updated value of  $RN_i[i]$ .)
- (b) When a site  $S_j$  receives this message, it sets  $RN_j[i]$  to  $\max(RN_j[i], sn)$ . If  $S_j$  has the idle token, then it sends the token to  $S_i$  if  $RN_j[i] = LN[i] + 1$ .

## Executing the critical section:

- (c) Site  $S_i$  executes the CS after it has received the token.

**Releasing the critical section:** Having finished the execution of the CS, site  $S_i$  takes the following actions:

- (d) It sets  $LN[i]$  element of the token array equal to  $RN_i[i]$ .
- (e) For every site  $S_j$  whose i.d. is not in the token queue, it appends its i.d. to the token queue if  $RN_i[j] = LN[j] + 1$ .
- (f) If the token queue is nonempty after the above update,  $S_i$  deletes the top site i.d. from the token queue and sends the token to the site indicated by the i.d.

---

**Algorithm 9.7** Suzuki–Kasami’s broadcast algorithm.

# Deadlock detection in distributed systems

- Deadlocks are a fundamental problem in distributed systems
- In distributed systems, a process may request resources in any order, which may not be known a priori and a process can request a resource while holding others.
- If the sequence of process resources is cyclic in such environments, deadlocks can occur.

Ktunotes.in

- A deadlock can be defined as a condition where a process is waiting for a resource that it already holds.
- Deadlocks can be dealt with using any one of the following three strategies: **deadlock prevention, deadlock avoidance, and deadlock detection.**
- **Deadlock prevention** is commonly achieved by either having a process release all its held resources or by suspending a process that holds the needed resource.

# Deadlock detection in distributed systems

- In the **deadlock avoidance** approach to distributed systems, a resource is granted to a process if the resulting global system is safe.
- **Deadlock detection** requires an examination of the \_\_\_\_\_ for the presence of a deadlock condition.
- To resolve the deadlock, we have to \_\_\_\_\_

## System model

Ktunotes.in

A distributed system consists of a set of processors that are connected by a communication network.

A distributed program is composed of a set of n asynchronous processes P1, P2, , Pi, , Pn that communicate by \_\_\_\_\_ over the communication network.

# Deadlock detection in distributed systems

- Without loss of generality we assume that each process is running on a different processor.
- The processors do not share memory and communicate solely by passing messages over the communication network.
- There is no global clock or shared memory.
- The communication medium may deliver messages out of order, messages may be lost, garbled, or duplicated due to timeout and retransmission, processors may fail, and communication links may go down.
- Deadlocks can occur in distributed systems.

www.IITunotes.in

# Deadlock detection in distributed systems

We make the following assumptions:

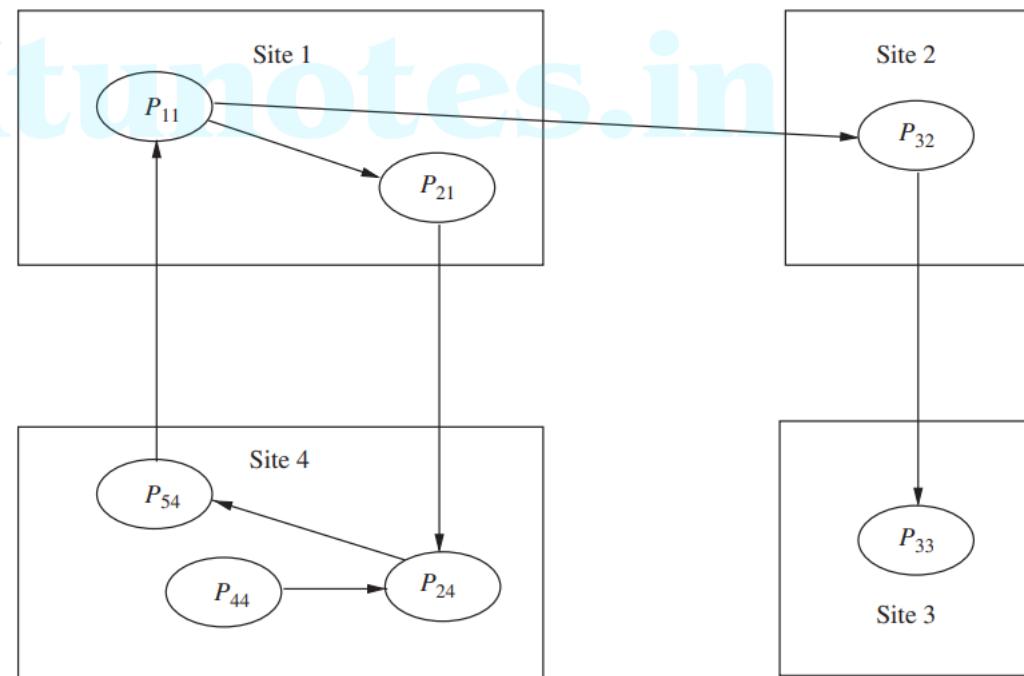
- The systems have finite number of processes.
- Processes are allowed to make only requests for resources.
- There is only one type of resource.
- A process can be in two states, running or blocked. In the running state (also called active state), a process has all the needed resources and is executing.
- In the blocked state, a process is waiting to acquire some resource.

# Deadlock detection in distributed systems

## Wait-for graph (WFG)

- In distributed systems, the state of the system can be modeled by directed graph, called a wait-for graph (WFG).
- In a WFG, nodes are processes and there is a directed edge from node P1 to node P2 if P1 is blocked and is waiting for P2 to release some resource.
- A system is deadlocked if and only if there exists a cycle in the WFG

**Figure 10.1** Example of a WFG.



# Deadlock detection in distributed systems

## Deadlock handling strategies

There are three strategies for handling deadlocks,

- deadlock prevention,
- deadlock avoidance,
- deadlock detection.

Handling of deadlocks becomes difficult in distributed systems because of the distributed nature and because every inter-site communication involves a finite and unpredictable delay.

**Deadlock prevention** is commonly achieved either by having a process acquire all the needed resources simultaneously before it begins executing or by pre-empting a process that holds the needed resource.

This approach is difficult to implement in distributed systems.

# Deadlock detection in distributed systems

In **deadlock avoidance** approach to distributed systems, a resource is granted to a process if the resulting global system state is safe.

Due to several problems, however, deadlock avoidance is impractical in distributed systems.

Deadlock detection requires an examination of the status of process- resource interactions for the presence of cyclic wait.

Deadlock detection in distributed systems seems to be the best approach to handle deadlocks in distributed systems.

# Deadlock detection in distributed systems

## Issues in deadlock detection

Deadlock handling using the approach of deadlock detection entails addressing two basic issues:

- detection of existing deadlocks
- resolution of detected deadlocks.

Ktunotes.in

### Detection of deadlocks

Detection of deadlocks involves addressing two issues:

Since, in distributed systems, a cycle or knot may involve several sites, the search for cycles greatly depends upon how the WFG of the system is represented across the system.

# Deadlock detection in distributed systems

Depending upon the way WFG information is maintained and the search for cycles is carried out

## Correctness criteria

A deadlock detection algorithm must

- 1. Progress (no undetected deadlocks)** : The algorithm must detect all existing deadlocks in a finite time.

x

- 2. Safety (no false deadlocks)** : The algorithm should not report deadlocks that do not exist (called **phantom deadlocks**).

In distributed systems where there is no global memory and there is no global clock, it is

As a result, sites may detect a cycle that never existed

# Deadlock detection in distributed systems

## Resolution of a detected deadlock

Deadlock resolution involves breaking existing wait-for dependencies between the processes to resolve the deadlock.

It involves **releasing resources** and assigning their resources to blocked processes so that they can resume execution

Ktunotes.in

# Deadlock detection in distributed systems

## Models of deadlocks

Distributed systems allow many kinds of resource requests.

A process might require a [redacted] for its execution

Models of deadlocks introduces a [redacted]

### 1. The single-resource model

The single-resource model is the [redacted] resource model in a distributed system, where a process can have at most one outstanding request for only one unit of a resource.

Since the [redacted] or the single resource model can be [redacted] the presence of [redacted] in the WFG shall indicate that there is a deadlock

# Deadlock detection in distributed systems

## 2. The AND model

- In the AND model, a process can request **more than one resource simultaneously** and the request is satisfied only after all the requested resources are granted to the process.
- The requested resources may exist **at different times**.
- The out degree of a node in the WFG for AND model can be more than 1.
- The presence of a **cycle** in the WFG indicates a deadlock in the AND model.

## 3. The OR model

In the OR model, a process can make a request for numerous resources simultaneously and the request is satisfied if any one of the requested resources is granted.

# Deadlock detection in distributed systems

The requested resources may exist at different locations.

If all requests in the WFG are OR requests, then the nodes are called **OR nodes**.

Presence of a **deadlock** in the OR model.

## 3. The AND-OR model

A generalization of the previous two models (OR model and AND model) is the AND-OR model.

In the AND-OR model, a request may **AND** or **OR** of and and or in the resource request.

For example, in the ANDOR model, a request for multiple resources can be of the form **AND** **OR**.

# Deadlock detection in distributed systems

## 4. The $\binom{p}{q}$ model

Another form of the AND-OR model is the  $(p q)$  model (called the P-out-of-Q model), which allows

Ktunotes.in

model lends itself to a much more compact formation of a request

# Deadlock detection in distributed systems

## 5. Unrestricted model

In the unrestricted model, no assumptions are made regarding the underlying structure of resource requests.

Only one assumption that the system is deterministic and hence it is the basis for deadlock detection.

Ktunotes.in

This model helps separate concerns: Communication and computation.

computations (e.g., message passing versus synchronous communication).